

FaceEmotionRecognition

SoniaMehra, SouvikPaul

31AUG 2020

ABSTRACT

In this Project we have used Convolutional Neural Network to recognise emotion of a face from its image. We have classified the emotions as 'Angry', 'Disgust', 'Fear', 'Happy', 'Sad', 'Surprise' and 'Neutral'. FER2013 data, consists of 35887 observations of images with emotion labels, is used. Firstly, the data is divided into training set, validation set and test set with ratio 8:1:1. Then from the training set we found out the number of observations corresponding to each label. Next, we trained the training data on different models and based on accuracy from the Validation set, we have chosen the best model and found its test accuracy. Lastly, we test our model with some real examples. Here, we have only provided the best of all the models which we have tried.

INTRODUCTION

In the era of Artificial Intelligence many unbelievable things become possible. Using a machine we can find out the person in front of you are telling lie or truth, just seeing a image of person Facebook can detect its name and details, hearing our voice Google understands and shows us the relevant websites etc. CNN is Deep Learning is such a tool to achieve perfection in image recognition. Significant advancement in image or object recognition were felt from 2011 to 2012. Although CNNs trained by backpropagation had been around for decades, and GPU implementations of NNs for years, including CNNs, fast implementations of CNNs with max-pooling on GPUs in the style of Cleeser and colleagues were needed to progress on computer vision. In 2011 this approach achieved for the first time superhuman performance in a visual pattern recognition contest. Also in 2011, it won the ICARD Chinese handwriting contest, and in May 2012, it won the ISBI image segmentation contest. Until 2011, CNNs did not play a major role at computer vision conferences, but in June 2012, a paper by Cleeser et al at the leading conference CVPR showed how max-pooling CNNs on GPU can drastically improve many vision benchmark records. In November 2012, Cleeser et al's system also won the KPCR contest on analysis of large medical images for cancer detection.

Here we also used CNN to recognise face emotion by seeing a picture of a person.

OBJECTIVE

Objective of the project is to develop a face emotion recogniser based on 30K samples.

DATA INFORMATION

This FER2013 data is collected from Kaggle. There are 35887 examples in the data and each example is 48 x 48 x 1 grayscale image. Data is divided into 80%, 10% and 10% respectively for train set, dev set and test set. Thus we get 28709 examples in the train set and 3589 examples in both the dev set and test set. This data was freely used in Kaggle Challenge and the winner got 71.01% test accuracy. The data consists of 3 columns and 5888 rows (with first row as heading). First column is for label, second column is for pixels (but all pixels are in a single excel cell separated by space for a single observation) and third column is for names of the emotion corresponding to label. In our training set there are 441 'angry' images (rows) and 7152 'angry' images. Let's do the analysis.

ANALYSIS

In the whole analysis various models are used. These models vary by hyperparameters such as number of convnet layers, number of fully connected layers, number of filters, size of filter, padding, strides, activation function. As there are seven categories, 'softmax' activation function is used. Final layer instead of 'softmax' (which is used for binary category). Keras is used in the analysis with Tensorflow at the backend.

Necessary modules to Import

```
In [33]: import tensorflow as tf
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import keras
import random
from sklearn.model_selection import train_test_split
import seaborn as sns
import os
import sys
import math

The function below is to import data in such a way that the pixel values for each observation can be elements of an array.
```

```
In [34]: def getData(filename):
    Y = []
    first = True
    for line in open(filename):
        if first:
            first = False
        else:
            row = line.split(',')
            Y.append(int(row[0]))
            X.append(int(row[1]) for p in row[1].split())
    X = np.array(X) / 255.0, np.array(Y)
    return X, Y
```

```
In [35]: from google.colab import drive
drive.mount('/content/drive')

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount('/content/drive', force_remount=True).
```

```
In [36]: X,Y=getData('/content/drive/My Drive/fer2013.csv')
```

Let's check the data.

```
In [37]: print(X)
print('Type of X is:',type(X))
print(Y)
print('Type of Y is:',type(Y))

[[0.2745988 0.5862329 0.2764709] ... 0.421568627 0.42745998 0.32156863]
[0.99588235 0.83137255 0.81176471 ... 0.34599894 0.43137255 0.59697843]

[0.96666667 0.96666667 0.9627451 ... 0.60392157 0.52156863 0.44313725]
[0.1764706 0.10980392 0.10980392 ... 0.1372549 0.11764706 0.10980392]
[0.674598 0.8599839 0.85496196 ... 0.74117647 0.78639216 0.78632529]]
Type of X is:
<class 'numpy.ndarray'>
[0 2 0 0 2]
Type of Y is:
<class 'numpy.ndarray'>
```

```
In [38]: X=X.reshape(35887,48,48,1)
print('Shape of X is:',X.shape)
print('Shape of Y is:',Y.shape)

Shape of X is: (35887, 48, 48, 1)
Shape of Y is: (35887,)
```

Here shape of X is (m,n_h,n_w,1) = (35887, 48, 48, 1) where m is the number of examples, n_h is number of pixels in height, n_w is number of pixels in width of an image and 1 due to single layer grayimage. Label Y is a vector of 35887 elements.

```
In [39]: X_train, X_test_val, y_train, y_test_val = train_test_split(X, y, test_size=0.2, random_state=8)
X_test, X_val, y_test, y_val = train_test_split(X_test_val, y_test_val, test_size=0.5, random_state=8, shuffle=False)
```

```
In [40]: print('Size of X_train:',X_train.shape[0])
print('Size of X_val:',X_val.shape[0])
print('Size of X_test:',X_test.shape[0])

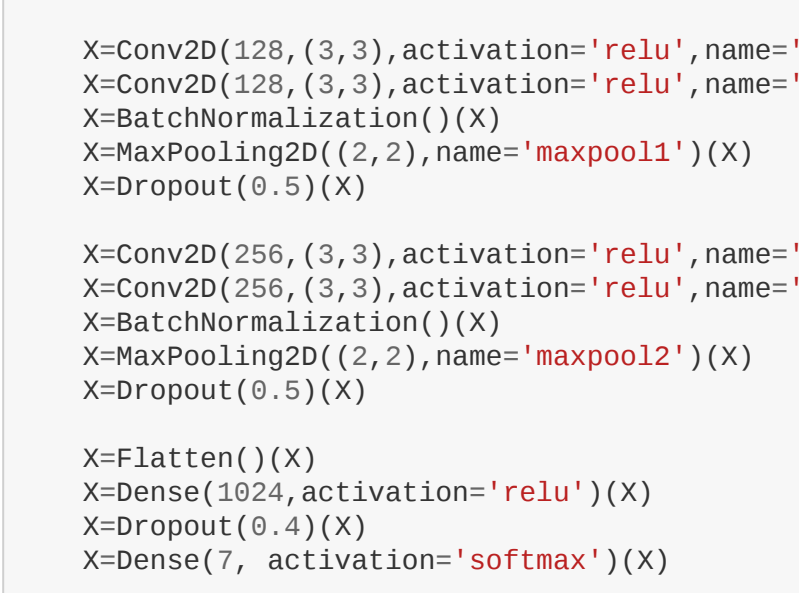
Size of X_train: 28709
Size of X_val: 3589
Size of X_test: 3589
```

Frequency distribution of each emotion in training set.

```
In [41]: emotion_map = {'Angry': 1, 'Disgust': 2, 'Fear': 3, 'Happy': 4, 'Sad': 5, 'Surprise': 6, 'Neutral': 0}
A=np.unique(y_train,return_counts=True)
emotion_counts=pd.DataFrame({'emotion':A,'number':B})
emotion_counts['emotion'] = emotion_counts['emotion'].map(emotion_map)
print(emotion_counts)
```

```
0      emotion    number
0      Angry      3952
1      Disgust      441
2      Fear       4115
3      Happy       7152
4      Sad        4516
5      Surprise    3285
6      Neutral     4990
```

```
In [42]: plt.figure(figsize=(8,4))
sns.barplot(emotion_counts.emotion, emotion_counts.number)
plt.title('Class distribution')
plt.ylabel('number', fontsize=12)
plt.xlabel('Emotions', fontsize=12)
plt.show()
```



Importing necessary modules for Keras.

```
In [43]: from keras.models import Sequential, Model
from keras.layers import Dense, Activation, Dropout, Flatten, ZeroPadding2D, Input
from keras.layers.convolutional import Conv2D
from keras.layers.convolutional import MaxPooling2D, AveragePooling2D
from keras.metrics import categorical_accuracy
from keras.models import model_from_json
from keras.callbacks import ModelCheckpoint
from keras.optimizers import *
from keras.preprocessing import image
import pydot
```

```
In [44]: random.seed(13)
```

Model

The structure of this model is shown below.

Conv2D0 -> Conv2D1 -> Conv2D2 -> BatchNormalization0 -> Maxpool0 -> Dropout -> Conv2D3 -> Conv2D4 -> BatchNormalization -> Maxpool1 -> Dropout -> Conv2D5 -> BatchNormalization -> Dropout -> FC0 -> Dropout -> FC1 -> Softmax

Note - (i) First argument in ZeroPadding2D is pad size.

(ii) First second argument in Conv2D is number of filters and filter size.

(iii) The argument in Activation is function relu is used.

(iv) Argument in MaxPooling2D is pool size.

(v) Flatten converts the Convnet into simple nn layer and Dense(7) fully connects this simple layer with another simple nn layer with 7 neurons

(vi) BatchNormalization is technique to normalize the input layer by re-centering and re-scaling

(vii) Dropout is a regularization technique. Dropout(p) randomly removes px100% neurons in a layer.

```
In [45]: def mymodel():
    input_shape=(48,48,1)
    X_input=Input(input_shape)

    X=Conv2D(64,(3,3),activation='relu',name='conv0')(X_input)
    X=Conv2D(64,(3,3),activation='relu',name='conv1')(X)
    X=Conv2D(64,(3,3),activation='relu',name='conv2')(X)
    X=BatchNormalization()(X)
    X=MaxPooling2D((2,2),name='maxpool0')(X)
    X=Dropout(0.5)(X)

    X=Conv2D(128,(3,3),activation='relu',name='conv3')(X)
    X=Conv2D(128,(3,3),activation='relu',name='conv4')(X)
    X=BatchNormalization()(X)
    X=MaxPooling2D((2,2),name='maxpool1')(X)
    X=Dropout(0.5)(X)

    X=Conv2D(256,(3,3),activation='relu',name='conv5')(X)
    X=Conv2D(256,(3,3),activation='relu',name='conv6')(X)
    X=BatchNormalization()(X)
    X=MaxPooling2D((2,2),name='maxpool2')(X)
    X=Dropout(0.5)(X)

    X=Flatten()(X)
    X=Dense(1024,activation='relu')(X)
    X=Dropout(0.4)(X)
    X=Dense(7,activation='softmax')(X)

    model=Model(inputs=X_input, outputs=X, name='mymodel')
    return model
```

```
In [46]: face_model=mymodel()
face_model.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['accuracy'])
y_train,y_val=train(np.arange(7)==y_train,None).astype(np.float32)
```

```
In [47]: y_val=(np.arange(7)==y_val).astype(np.float32)
```

```
In [68]: history=face_model.fit(x=X_train, y=y_train, epochs=50, batch_size=64, verbose=1, validation_data=(X_val,y_val))
```

```
Epoch 1/50 ..... ETA: 11s - loss: 0.5333 - accuracy: 0.8125WARNING:
tensorflow.CallBacks method on 'train_batch_end' is slow compared to the batch time (batch 11 ne: 0.0038 vs on_tty1 batch time: 0.0206). Check your callbacks.
449/449 [=====] - val_loss: 0.6448 - accuracy: 0.7571 - val_loss: 1.0789 - val_accuracy: 0.6590
Epoch 2/50 .....
449/449 [=====] - val_loss: 1.0932 - val_accuracy: 0.6590
Epoch 3/50 .....
449/449 [=====] - val_loss: 1.2947 - val_accuracy: 0.6514
Epoch 4/50 .....
449/449 [=====] - val_loss: 1.1123 - val_accuracy: 0.6595
Epoch 5/50 .....
449/449 [=====] - val_loss: 1.0685 - val_accuracy: 0.6601
Epoch 6/50 .....
449/449 [=====] - val_loss: 1.0802 - val_accuracy: 0.6604
Epoch 7/50 .....
449/449 [=====] - val_loss: 1.1135 - val_accuracy: 0.6525
Epoch 8/50 .....
449/449 [=====] - val_loss: 1.1178 - val_accuracy: 0.6601
Epoch 9/50 .....
449/449 [=====] - val_loss: 1.1154 - val_accuracy: 0.6623
Epoch 10/50 .....
449/449 [=====] - val_loss: 1.0807 - val_accuracy: 0.6641
Epoch 11/50 .....
449/449 [=====] - val_loss: 1.0851 - val_accuracy: 0.6590
Epoch 12/50 .....
449/449 [=====] - val_loss: 1.1757 - val_accuracy: 0.6617
Epoch 13/50 .....
449/449 [=====] - val_loss: 1.1139 - val_accuracy: 0.6573
Epoch 14/50 .....
449/449 [=====] - val_loss: 1.1491 - val_accuracy: 0.6498
Epoch 15/50 .....
449/449 [=====] - val_loss: 1.1447 - val_accuracy: 0.6567
Epoch 16/50 .....
449/449 [=====] - val_loss: 1.1037 - val_accuracy: 0.6634
Epoch 17/50 .....
449/449 [=====] - val_loss: 1.1026 - val_accuracy: 0.6634
Epoch 18/50 .....
449/449 [=====] - val_loss: 1.1169 - val_accuracy: 0.6484
Epoch 19/50 .....
449/449 [=====] - val_loss: 1.1055 - val_accuracy: 0.6525
Epoch 20/50 .....
449/449 [=====] - val_loss: 1.1427 - val_accuracy: 0.6434
Epoch 21/50 .....
449/449 [=====] - val_loss: 1.1674 - val_accuracy: 0.6567
Epoch 22/50 .....
449/449 [=====] - val_loss: 1.1699 - val_accuracy: 0.6523
Epoch 23/50 .....
449/449 [=====] - val_loss: 1.1087 - val_accuracy: 0.6590
Epoch 24/50 .....
449/449 [=====] - val_loss: 1.1074 - val_accuracy: 0.6490
Epoch 25/50 .....
449/449 [=====] - val_loss: 1.1059 - val_accuracy: 0.6584
Epoch 26/50 .....
449/449 [=====] - val_loss: 1.1169 - val_accuracy: 0.6578
Epoch 27/50 .....
449/449 [=====] - val_loss: 1.1866 - val_accuracy: 0.6528
Epoch 28/50 .....
449/449 [=====] - val_loss: 1.1049 - val_accuracy: 0.6567
Epoch 29/50 .....
449/449 [=====] - val_loss: 1.1485 - val_accuracy: 0.6562
Epoch 30/50 .....
449/449 [=====] - val_loss: 1.1464 - val_accuracy: 0.6634
Epoch 31/50 .....
449/449 [=====] - val_loss: 1.1464 - val_accuracy: 0.6586
Epoch 32/50 .....
449/449 [=====] - val_loss: 1.1464 - val_accuracy: 0.6586
Epoch 33/50 .....
449/449 [=====] - val_loss: 1.1249 - val_accuracy: 0.6562
Epoch 34/50 .....
449/449 [=====] - val_loss: 1.1393 - val_accuracy: 0.6648
Epoch 35/50 .....
449/449 [=====] - val_loss: 1.1623 - val_accuracy: 0.6634
Epoch 36/50 .....
449/449 [=====] - val_loss: 1.1985 - val_accuracy: 0.6428
Epoch 37/50 .....
449/449 [=====] - val_loss: 1.1912 - val_accuracy: 0.6539
Epoch 38/50 .....
449/449 [=====] - val_loss: 1.2578 - val_accuracy: 0.6584
Epoch 39/50 .....
449/449 [=====] - val_loss: 1.2297 - val_accuracy: 0.6547
Epoch 40/50 .....
449/449 [=====] - val_loss: 1.1904 - val_accuracy: 0.6587
Epoch 41/50 .....
449/449 [=====] - val_loss: 1.2251 - val_accuracy: 0.6551
Epoch 42/50 .....
449/449 [=====] - val_loss: 1.1816 - val_accuracy: 0.6578
Epoch 43/50 .....
449/449 [=====] - val_loss: 1.1813 - val_accuracy: 0.6581
Epoch 44/50 .....
449/449 [=====] - val_loss: 1.1339 - val_accuracy: 0.6583
Epoch 45/50 .....
449/449 [=====] - val_loss: 1.1933 - val_accuracy: 0.6539
Epoch 46/50 .....
449/449 [=====] - val_loss: 1.1759 - val_accuracy: 0.6598
Epoch 47/50 .....
449/449 [=====] - val_loss: 1.1697 - val_accuracy: 0.6567
Epoch 48/50 .....
449/449 [=====] - val_loss: 1.1895 - val_accuracy: 0.6542
Epoch 49/50 .....
449/449 [=====] - val_loss: 1.1867 - val_accuracy: 0.6525
Epoch 50/50 .....
449/449 [=====] - val_loss: 1.1967 - val_accuracy: 0.6525
Epoch 51/50 .....
449/449 [=====] - val_loss: 1.1904 - val_accuracy: 0.6587
Epoch 52/50 .....
449/449 [=====] - val_loss: 1.2251 - val_accuracy: 0.6551
Epoch 53/50 .....
449/449 [=====] - val_loss: 1.1816 - val_accuracy: 0.6578
Epoch 54/50 .....
449/449 [=====] - val_loss: 1.1813 - val_accuracy: 0.6581
Epoch 55/50 .....
449/449 [=====] - val_loss: 1.1339 - val_accuracy: 0.6583
Epoch 56/50 .....
449/449 [=====] - val_loss: 1.1933 - val_accuracy: 0.6539
Epoch 57/50 .....
449/449 [=====] - val_loss: 1.1759 - val_accuracy: 0.6598
Epoch 58/50 .....
449/449 [=====] - val_loss: 1.1697 - val_accuracy: 0.6567
Epoch 59/50 .....
449/449 [=====] - val_loss: 1.1895 - val_accuracy: 0.6542
Epoch 60/50 .....
449/449 [=====] - val_loss: 1.1867 - val_accuracy: 0.6525
Epoch 61/50 .....
449/449 [=====] - val_loss: 1.1967 - val_accuracy: 0.6525
Epoch 62/50 .....
449/449 [=====] - val_loss: 1.1904 - val_accuracy: 0.6587
Epoch 63/50 .....
449/449 [=====] - val_loss: 1.2251 - val_accuracy: 0.6551
Epoch 64/50 .....
449/449 [=====] - val_loss: 1.1816 - val_accuracy: 0.6578
Epoch 65/50 .....
449/449 [=====] - val_loss: 1.1813 - val_accuracy: 0.6581
Epoch 66/50 .....
449/449 [=====] - val_loss: 1.1339 - val_accuracy: 0.6583
Epoch 67/50 .....
449/449 [=====] - val_loss: 1.1933 - val_accuracy: 0.6539
Epoch 68/50 .....
449/449 [=====] - val_loss: 1.1759 - val_accuracy: 0.6598
Epoch 69/50 .....
449/449 [=====] - val_loss: 1.1697 - val_accuracy: 0.6567
Epoch 70/50 .....
449/449 [=====] - val_loss: 1.1895 - val_accuracy: 0.6542
Epoch 71/50 .....
449/449 [=====] - val_loss: 1.1867 - val_accuracy: 0.6525
Epoch 72/50 .....
449/449 [=====] - val_loss: 1.1967 - val_accuracy: 0.6525
Epoch 73/50 .....
449/449 [=====] - val_loss: 1.1904 - val_accuracy: 0.6587
Epoch 74/50 .....
449/449 [=====] - val_loss: 1.2251 - val_accuracy: 0.6551
Epoch 75/50 .....
449/449 [=====] - val_loss: 1.1816 - val_accuracy: 0.6578
Epoch 76/50 .....
449/449 [=====] - val_loss: 1.1813 - val_accuracy: 0.6581
Epoch 77/50 .....
449/449 [=====] - val_loss: 1.1339 - val_accuracy: 0.6583
Epoch 78/50 .....
449/449 [=====] - val_loss: 1.1933 - val_accuracy: 0.6539
Epoch 79/50 .....
449/449 [=====] - val_loss: 1.1759 - val_accuracy: 0.6598
Epoch 80/50 .....
449/449 [=====] - val_loss: 1.1697 - val_accuracy: 0.6567
Epoch 81/50 .....
449/449 [=====] - val_loss: 1.1895 - val_accuracy: 0.6542
Epoch 82/50 .....
449/449 [=====] - val_loss: 1.1867 - val_accuracy: 0.6525
Epoch 83/50 .....
449/449 [=====] - val_loss: 1.1967 - val_accuracy: 0.6525
Epoch 84/50 .....
449/449 [=====] - val_loss: 1.1904 - val_accuracy: 0.6587
Epoch 85/50 .....
449/449 [=====] - val_loss: 1.2251 - val_accuracy: 0.6551
Epoch 86/50 .....
449/449 [=====] - val_loss: 1.1816 - val_accuracy: 0.6578
Epoch 87/50 .....
449/449 [=====] - val_loss: 1.1813 - val_accuracy: 0.6581
Epoch 88/50 .....
449/449 [=====] - val_loss: 1.1339 - val_accuracy: 0.6583
Epoch 89/50 .....
449/449 [=====] - val_loss: 1.1933 - val_accuracy: 0.6539
Epoch 90/50 .....
449/449 [=====] - val_loss: 1.1759 - val_accuracy: 0.6598
Epoch 91/50 .....
449/449 [=====] - val_loss: 1.1697 - val_accuracy: 0.6567
Epoch 92/50 .....
449/449 [=====] - val_loss: 1.1895 - val_accuracy: 0.6542
Epoch 93/50 .....
449/449 [=====] - val_loss: 1.1867 - val_accuracy: 0.6525
Epoch 94/50 .....
449/449 [=====] - val_loss: 1.1967 - val_accuracy: 0.6525
Epoch 95/50 .....
449/449 [=====] - val_loss: 1.1904 - val_accuracy: 0.6587
Epoch 96/50 .....
449/449 [=====] - val_loss: 1.2251 - val_accuracy: 0.6551
Epoch 97/50 .....
449/449 [=====] - val_loss: 1.1816 - val_accuracy: 0.6578
Epoch 98/50 .....
449/449 [=====] - val_loss: 1.1813 - val_accuracy: 0.6581
Epoch 99/50 .....
449/449 [=====] - val_loss: 1.1339 - val_accuracy: 0.6583
Epoch 100/50 .....
449/449 [=====] - val_loss: 1.1933 - val_accuracy: 0.6539
Epoch 101/50 .....
449/449 [=====] - val_loss: 1.1759 - val_accuracy: 0.6598
Epoch 102/50 .....
449/449 [=====] - val_loss: 1.1697 - val_accuracy: 0.6567
Epoch 103/50 .....
449/449 [=====] - val_loss: 1.1895 - val_accuracy: 0.6542
Epoch 104/50 .....
449/449 [=====] - val_loss: 1.1867 - val_accuracy: 0.6525
Epoch 105/50 .....
449/449 [=====] - val_loss: 1.1967 - val_accuracy: 0.6525
Epoch 106/50 .....
449/449 [=====] - val_loss: 1.1904 - val_accuracy: 0.6587
Epoch 107/50 .....
449/449 [=====] - val_loss: 1.2251 - val_accuracy: 0.6551
Epoch 108/50 .....
449/449 [=====] - val_loss: 1.1816 - val_accuracy: 0.6578
Epoch 109/50 .....
449/449 [=====] - val_loss: 1.1813 - val_accuracy: 0.6581
Epoch 110/50 .....
449/449 [=====] - val_loss: 1.1339 - val_accuracy: 0.6583
Epoch 111/50 .....
449/449 [=====] - val_loss: 1.1933 - val_accuracy: 0.6539
Epoch 112/50 .....
449/449 [=====] - val_loss: 1.1759 - val_accuracy: 0.6598
Epoch 113/50 .....
449/449 [=====] - val_loss: 1.1697 - val_accuracy: 0.6567
Epoch 114/50 .....
449/449 [=====] - val_loss: 1.1895 - val_accuracy: 0.6542
Epoch 115/50 .....
449/449 [=====] - val_loss: 1.1867 - val_accuracy: 0.6525
Epoch 116/50 .....
449/449 [=====] - val_loss: 1.1967 - val_accuracy: 0.6525
Epoch 117/50 .....
449/449 [=====] - val_loss: 1.1904 - val_accuracy: 0.6587
Epoch 118/50 .....
449/449 [=====] - val_loss: 1.2251 - val_accuracy: 0.6551
Epoch 119/50 .....
449/449 [=====] - val_loss: 1.1816 - val_accuracy: 0.6578
Epoch 120/50 .....
449/449 [=====] - val_loss: 1.1813 - val_accuracy: 0.6581
Epoch 121/50 .....
449/449 [=====] - val_loss: 1.1339 - val_accuracy: 0.6583
Epoch 122/50 .....
449/449 [=====] - val_loss: 1.1933 - val_accuracy: 0.6539
Epoch 123/50 .....
449/449 [=====] - val_loss: 1.1759 - val_accuracy: 0.6598
Epoch 124/50 .....
449/449 [=====] - val_loss: 1.1697 - val_accuracy: 0.6567
Epoch 125/50 .....
449/449 [=====] - val_loss: 1.1895 - val_accuracy: 0.6542
Epoch 126/50 .....
449/449 [=====] - val_loss: 1.1867 - val_accuracy: 0.6525
Epoch 127/50 .....
449/449 [=====] - val_loss: 1.1967 - val_accuracy: 0.6525
Epoch 128/50 .....
449/449 [=====] - val_loss: 1.1904 - val_accuracy: 0.6587
Epoch 129/50 .....
449/449 [=====] - val_loss: 1.2251 - val_accuracy: 0.6551
Epoch 130/50 .....
449/449 [=====] - val_loss: 1.1816 - val_accuracy: 0.6578
Epoch 131/50 .....
449/449 [=====] - val_loss: 1.1813 - val_accuracy: 0.6581
Epoch 132/50 .....
449/449 [=====] - val_loss: 1.1339 - val_accuracy: 0.6583
Epoch 133/50 .....
449/449 [=====] - val_loss: 1.1933 - val_accuracy: 0.6539
Epoch 134/50 .....
449/449 [=====] - val_loss: 1.1759 - val_accuracy: 0.6598
Epoch 135/50 .....
449/449 [=====] - val_loss: 1.1697 - val_accuracy: 0.6567
Epoch 136/50 .....
449/449 [=====] - val_loss: 1.1895 - val_accuracy: 0.6542
Epoch 137/50 .....
449/449 [=====] - val_loss: 1.1867 - val_accuracy: 0.6525
Epoch 138/50 .....
449/449 [=====] - val_loss: 1.1967 - val_accuracy: 0.6525
Epoch 139/50 .....
449/449 [=====] - val_loss: 1.1904 - val_accuracy: 0.6587
Epoch 140/50 .....
449/449 [=====] - val_loss: 1.2251 - val_accuracy: 0.6551
Epoch 141/50 .....
449/449 [=====] - val_loss: 1.1816 - val_accuracy: 0.6578
Epoch 142/50 .....
449/449 [=====] - val_loss: 1.1813 - val_accuracy: 0.6581
Epoch 143/50 .....
449/449 [=====] - val_loss: 1.1339 - val_accuracy: 0.6583
Epoch 144/50 .....
449/449 [=====] - val_loss: 1.1933 - val_accuracy: 0.6539
Epoch 145/50 .....
449/449 [=====] - val_loss: 1.1759 - val_accuracy: 0.6598
Epoch 146/50 .....
449/449 [=====] - val_loss: 1.1697 - val_accuracy: 0.6567
Epoch 147/50 .....
449/449 [=====] - val_loss: 1.1895 - val_accuracy: 0.6542
Epoch 148/50 .....
449/449 [=====] - val_loss: 1.1867 - val_accuracy: 0.6525
Epoch 149/50 .....
449/449 [=====] - val_loss: 1.1967 - val_accuracy: 0.6525
Epoch 150/50 .....
449/449 [=====] - val_loss: 1.1904 - val_accuracy: 0.6587
Epoch 151/50 .....
449/449 [=====] - val_loss: 1.2251 - val_accuracy: 0.6551
Epoch 152/50 .....
449/449 [=====] - val_loss: 1.1816 - val_accuracy: 0.6578
Epoch 153/50 .....
449/449 [=====] - val_loss: 1.1813 - val_accuracy: 0.6581
Epoch 154/50 .....
449/449 [=====] - val_loss: 1.1339 - val_accuracy: 0.6583
Epoch 155/50 .....
449/449 [=====] - val_loss: 1.1933 - val_accuracy: 0.6539
Epoch 156/50 .....
449/449 [=====] - val_loss: 1.1759 - val_accuracy: 0.6598
Epoch 157/50 .....
449/449 [=====] - val_loss: 1.1697 - val_accuracy: 0.6567
Epoch 158/50 .....
449/449 [=====] - val_loss: 1.1895 - val_accuracy: 0.6542
Epoch 159/50 .....
449/449 [=====] - val_loss: 1.1867 - val_accuracy: 0.6525
Epoch 160/50 .....
449/449 [=====] - val_loss: 1.1967 - val_accuracy: 0.6525
Epoch 161/50 .....
449/449 [=====] - val_loss: 1.1904 - val_accuracy: 0.6587
Epoch 162/50 .....
449/449 [=====] - val_loss: 1.2251 - val_accuracy: 0.6551
Epoch 163/50 .....
449/449 [=====] - val_loss: 1.1816 - val_accuracy: 0.6578
Epoch 164/50 .....
449/449 [=====] - val_loss: 1.1813 - val_accuracy: 0.6581
Epoch 165/50 .....
449/449 [=====] - val_loss: 1.1339 - val_accuracy: 0.6583
Epoch 166/50 .....
449/449 [=====] - val_loss: 1.1933 - val_accuracy: 0.6539
Epoch 167/50 .....
449/449 [=====] - val_loss: 1.1759 - val_accuracy: 0.6598
Epoch 168/50 .....
449/449 [=====] - val_loss: 1.1697 - val_accuracy: 0.6567
Epoch 169/50 .....
449/449 [=====] - val_loss: 1.1895 - val_accuracy: 0.6542
Epoch 170/50 .....
449/449 [=====] - val_loss: 1.1867 - val_accuracy: 0.6525
Epoch 171/50 .....
449/449 [=====] - val_loss: 1.1967 - val_accuracy: 0.6525
Epoch 172/50 .....
449/449 [=====] - val_loss: 1.1904 - val_accuracy: 0.6587
Epoch 173/50 .....
449/449 [=====] - val_loss: 1.2251 - val_accuracy: 0.6551
Epoch 174/50 .....
449/449 [=====] - val_loss: 1.1816 - val_accuracy: 0.6578
Epoch 175/50 .....
449/449 [=====] - val_loss: 1.1813 - val_accuracy: 0.6581
Epoch 176/50 .....
449/449 [=====] - val_loss: 1.1339 - val_accuracy: 0.6583
Epoch 177/50 .....
449/449 [=====] - val_loss: 1.1933 - val_accuracy: 0.6539
Epoch 178/50 .....
449/449 [=====] - val_loss: 1.1759 - val_accuracy: 0.6598
Epoch 179/50 .....
449/449 [=====] - val_loss: 1.1697 - val_accuracy: 0.6567
Epoch 180/50 .....
449/449 [
```